

**CS528**  
**C++ opt**  
**&**  
**Data Access Optimization**

A Sahu  
Dept of CSE, IIT Guwahati

# Outline

- C++ Optimization
- **Data Access Optimization**
- **Roof line Model**

# C++ Optimizations

- Temporaries
- Dynamic Memory Management
- Loop Kernel and Iterators

# C++ Opt: Temporaries

- C++: operator overloading uses

```
class vec3d{
    double x,y,z;
public: vec3d( double _x=0.0, _y=0.0, _z=0.0):x(_x),y(-y),z(_z){}
    vec3d operator+(const vec3d &oth){
        vec3d tmp; tmp.x=x+oth.x; ...for y, and z
        return tmp
    }
    vec3d operator*(double s, const vec3d &v){
        vec3d tmp(s*v.x, s*v.y,s*v.z); return tmp;}
}
main() {
    vec3d a, b(2,2), c(3); double u=1.0,v=2.0;
    a=u*b + v*c;
}
```

# C++ Opt: Temporaries

- C++: operator overloading uses
- In this prev statements
  - Constructor get called for a,b,c
  - Operator\*, constructor for tmp, destructor for tmp
  - Operator\*, constructor for tmp, destructor for tmp
  - Operator+, constructor for tmp, destructor for tmp
  - Copy constrtor called with tmp
- **Simply we could have write**
  - **`a.x=u*b.x+v*c.x; a.y=u*b.y+v*c.y; a.z=u*b.z+v*c.z;`**

# C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
    vector<double> v(Len);  
    if(rand() > Th * RAND_MAX) {  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



This creation is  
Costly

# C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
  
    if(rand() > Th * RAND_MAX) {  
        vector<double> v(Len);  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



**This creation is  
Costly, so make it  
Lazy**

- **Lazy construction** : if the probability of requirement is low
  - Post pone the construction if the condition become true

# C++ Opt: Dynamic Memory Management

```
void func(double Th, int Len) {  
    static vector<double> v(LargeLen);  
    if(rand() > Th * RAND_MAX) {  
        v = obtain_data(Len);  
        sort(v.begin(), v.end());  
        process_data(v);  
    }  
}
```



**One time  
construction for  
all calls**

- **Static Construction** : if the probability of requirement is high or always required
  - one time Construction : for all call/invocation
  - Take sufficient largeLen



# C++ Opt: Loop Kernel and Iterators

- Runtime of scientific application dominated by loops or loops nest
- Compiler ability to optimize loops is pivotal for getting performance
- Operator overloading and template may hinders good loop optimization

# C++ Opt: Loop Kernel and Iterators

- Non-SIMDized code: operator[] called twice for a and b, compiler refuse to SIMDize

```
template<class T>
T Sprod(const vector<T> &A,
        const vector<T> &B) {

    T result=T(0);
    int s=A.size();
    for(int i=0;i<s;i++)
        result += A[i]*B[i]; //Access
    return result;
}
```

# C++ Opt: Loop Kernel and Iterators

- SIMDized

```
template<class T>
T Sprod(const vector<T> &A,
        const vector<T> &B) {
    vector<T>::const_iterator
        iA=A.begin(), iB=B.begin();
    T result=T(0);
    int s=A.size();
    for(int i=0;i<s;i++)
        result += iA[i]*iB[i]; //Access
    return result;
}
```

# Data Access Optimization

# DAO

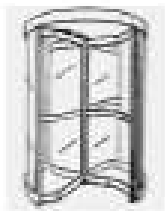
- Data Access Optimization
  - Roofline Model
  - Caching optimization
  - App classification based DA:  $N/N$ ,  $N^2/N^2$ ,  $N^3/N^2$
- *[Ref: Hager Book, PDF uploaded to Website]*

# Performance of System: Modeling Customer Dispatch in a Bank

Resolving door

Throughput:

$b_s$  [customer/sec]



Processing  
Capability:

$P_{peak}$  [task/sec]

Intensity:

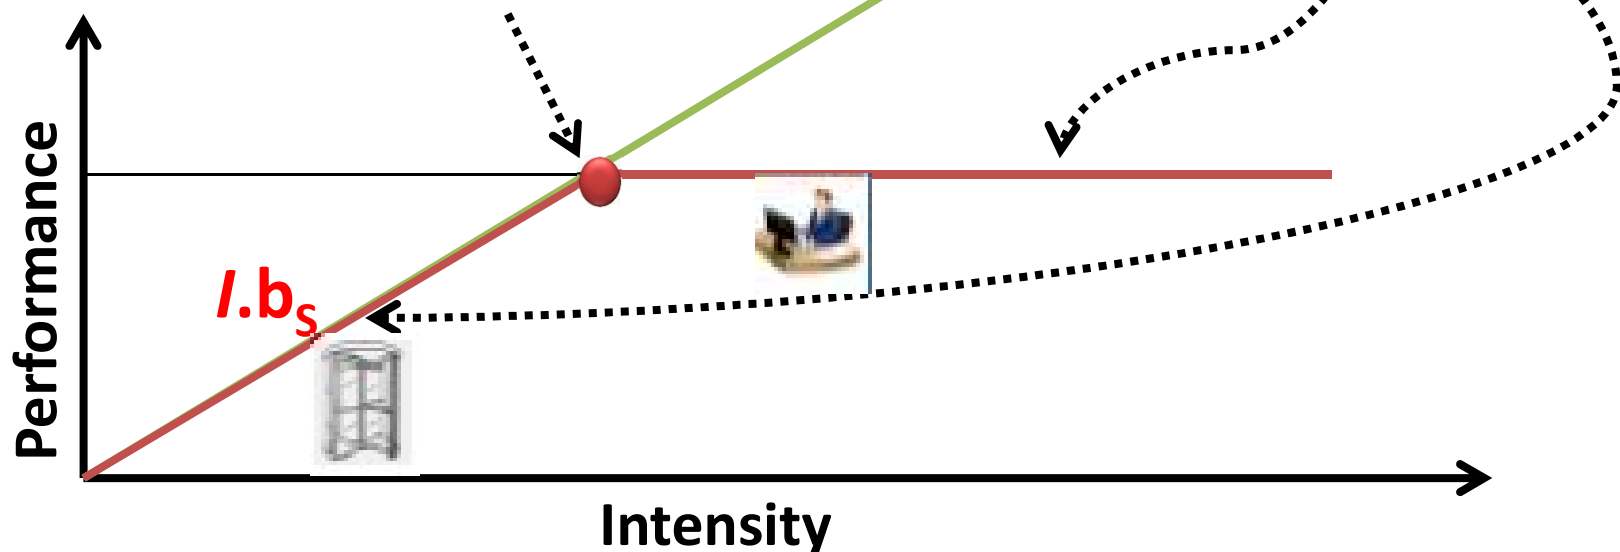
$I$  [task/customer]

# Modeling Customer Dispatch in a Bank

- How fast can tasks be processed?  $P[\text{tasks/sec}]$
- The bottleneck is either
  - The service desks (peak. tasks/sec):  $P_{\text{peak}}$
  - The revolving door (max. customers/sec):  $I \cdot b_s$
- Performance  $P = \min(P_{\text{peak}}, I \cdot b_s)$
- This is the “Roofline Model”
  - High intensity:  $P$  limited by “execution”
  - Low intensity:  $P$  limited by “bottleneck”

# Modeling Customer Dispatch in a Bank

- Performance  $P = \min(P_{\text{peak}}, I \cdot b_s)$
- This is the “Roofline Model”
  - High intensity: P limited by “execution”
  - Low intensity: P limited by “bottleneck”
  - “Knee” at  $P_{\text{peak}} = I \cdot b_s$ : Best use of resources



- Roofline is an “optimistic” model



# The Roofline Model

- $P_{\max}$  = Peak performance of the machine
- $I$  = Computational intensity (“work” per byte transferred) over the slowest data path utilized (“the bottleneck”)
- $b_s$  = Applicable peak bandwidth of the slowest data path utilized

Expected performance:

$$P = \min(P_{\text{peak}}, I \cdot b_s)$$

[F/B]

[B/s]

# Apply Roof line to Machine and Code

- Machine Parameter 1 :  $P_{\text{peak}} [\text{F/s}] = 4 \text{ G F/s}$
- Machine Parameter 2 :  $b_s [\text{B/s}] = 10 \text{ G B/s}$
- Application Properties:  $I [\text{F/B}] = 2\text{F}/8\text{B} = 0.25\text{F/B}$   
**for(i=0;i<N;i++) s=s+a[i]\*a[i]; // double s, a[]**
- Performance =  $P = \min(P_{\text{peak}}, I * b_s)$   
 $= \min(4 \text{ GF/s}, 0.25 \text{ F/B} * 10 \text{ G.B/s})$   
 $= \min(4 \text{ GF/s}, 2.5 \text{ GF/s})$   
 $= 2.5 \text{ G F/s}$

# The Refine Roofline Model

- $P_{\max}$  = Peak performance of a loop assuming that data comes from L1 cache (**is not necessarily  $P_{\text{peak}}$** )
- $I$  = Computational intensity (“work” per byte transferred) over the slowest data path utilized (“the bottleneck”),
- **Code Balance  $B_c = I^{-1}$  = in Byte per Flop**
- $b_s$  = Applicable peak bandwidth of the slowest data path utilized

Expected performance:


$$P = \min(P_{\max}, I \cdot b_s)$$

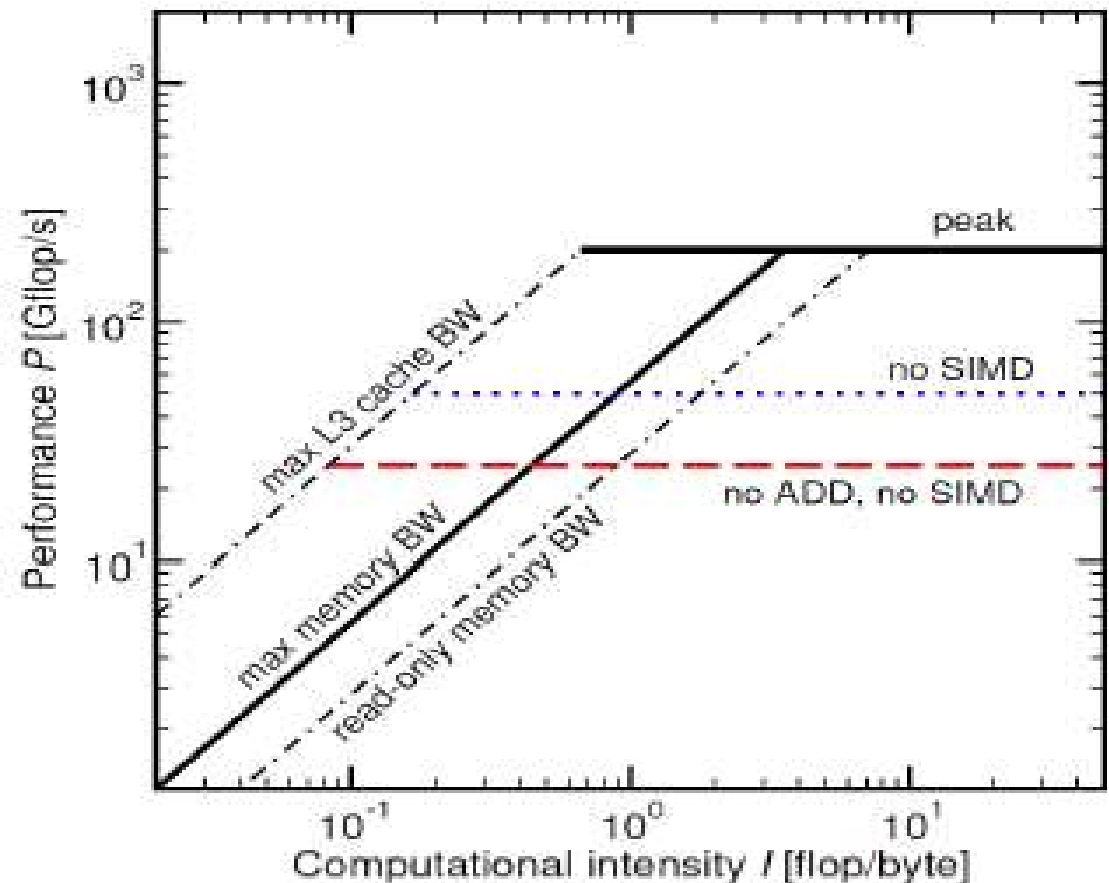
$$P = \min(P_{\max}, b_s / B_c)$$

# Factors to consider in Roofline Model

- BW Bound : may be simple
  - Accurate traffic calculation (write allocate, stride,..)
  - Practical not equal to theoretical BW limits
  - Saturation effects -> consider full socket only
- Core Bound : may be complex
  - Multiple bottlenecks: LD/ST, arithmetic, pipeline, SIMD, execution port
  - Limit is linear in # of cores

# Refine RFL model: Graphical Representation

- Multiple ceiling may apply
  - Diff BW/Data paths  
-> Diff inclined ceilings
  - Different  $P_{\max}$  ->  
Diff flat ceilings



**$P_{\max}$  comes from code analysis :**  
**with/without SIMD, add other FUs**

# Apply Roof line to Haswell Core to Triad Code

- Achievable Max Performance 1:  $P_{\max} [\text{F/s}] = 12.27 \text{ G F/s}$
- Machine Parameter 2 :  $b_s [\text{B/s}] = 50 \text{ G B/s}$
- Application Properties:  $I [\text{F/B}] = 2F/40B = 0.05 \text{ F/B}$   
**for(i=0;i<N;i++) a[i]=b[i]+c[i]\*d[i]; // double a,b,c,d**
- Performance =  $P = \min(P_{\max}, I * b_s)$   
 $= \min(12.27 \text{ GF/s}, 0.05 \text{ F/B} * 50 \text{ G.B/s})$   
 $= \min(12.27 \text{ GF/s}, 2.5 \text{ GF/s})$   
 $= 2.5 \text{ G F/s}$

# Code Balance/Intensity Examples

```
for (i=0; i<N; i++) //Copy  
    a[i]=b[i];
```

$$B_c = 24B/0F = NA$$

```
for (i=0; i<N; i++) //Scale  
    a[i]=s*b[i];
```

$$B_c = 24B/1F = 24 \text{ B/F}$$

```
for (i=0; i<N; i++) //Add  
    a[i]= b[i]*c[i];
```

$$B_c = 32B/1F = 32 \text{ B/F}$$

`A[i]=B[i];` //Require reading of A[i], B[i] and writing to A[i]  
incase of L1 cache access; so 2 read and one write

# Code Balance/Intensity Examples

```
double a[N], b[N], c[N], d[N];  
for (i=0; i<N; i++)  
    a[i] = a[i] + b[i];
```

$$B_c = 24B/1F = 24 \text{ B/F}$$
$$I = 0.042 F/B$$

```
for (i=0; i<N; i++) //Triad  
    a[i] = a[i] + s*b[i];
```

$$B_c = 24B/2F = 12 \text{ B/F}$$
$$I = 0.083 F/B$$

```
for (i=0; i<N; i++) //float a[]  
    s = s + a[i]*a[i];
```

$$B_c = 4B/2F = 2 \text{ B/F}$$
$$I = 0.5 F/B$$

```
for (i=0; i<N; i++) //float a[], b[]  
    s = s + a[i]*b[i];
```

$$B_c = 8B/2F = 4 \text{ B/F}$$
$$I = 0.25 F/B$$