

CS528

OpenMP and MPI

A Sahu

Dept of CSE, IIT Guwahati

Outline

- OpenMP
- MPI

OpenMP

OpenMP

- Compiler directive: Automatic parallelization
- Auto generate thread and get synchronized

```
#include <openmp.h>
main() {
    #pragma omp parallel
    #pragma omp for schedule(static)
    {
        for (int i=0; i<N; i++) {
            a[i]=b[i]+c[i];
        }
    }
}
```

```
$ gcc -fopenmp test.c
```

```
$ export OMP_NUM_THREADS=4
```

```
$/a.out
```

OpenMP: Parallelism

Sequential code

```
for (int i=0; i<N; i++)  
    a[i]=b[i]+c[i];
```

OpenMP: Parallelism

(Semi) manual parallel

```
#pragma omp parallel
{
    int id =omp_get_thread_num();
    int Nthr=omp_get_num_threads();
    int istart = id*N/Nthr
    int iend= (id+1)*N/Nthr;
    for (int i=istart;i<iend;i++) {
        a[i]=b[i]+c[i];
    }
}
```

OpenMP: Parallelism

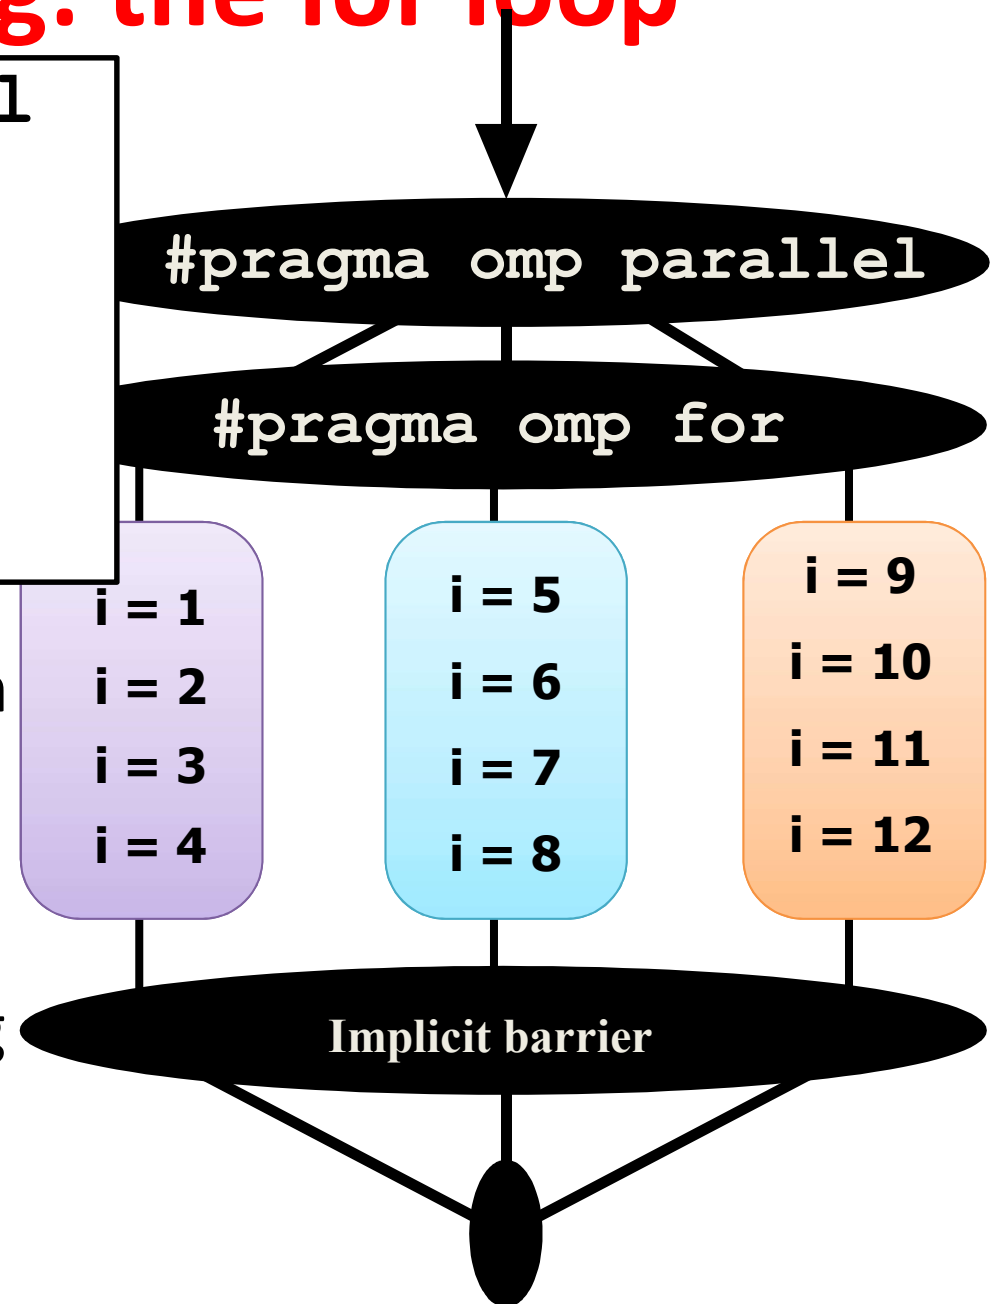
Auto parallel for loop

```
#pragma omp parallel
#pragma omp for schedule(static)
{
    for (int i=0; i<N; i++) {
        a[i]=b[i]+c[i];
    }
}
```

Work-sharing: the for loop


```
#pragma omp parallel
#pragma omp for
{
    for (i=1; i<13; i++)
        c[i]=a[i]+b[i];
}
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct



OpenMP Fork-and-Join model

```
printf("begin\n");  
N = 1000;  
  
#pragma omp parallel for  
for (i=0; i<N; i++)  
    A[i] = B[i] + C[i];  
  
M = 500;  
  
#pragma omp parallel for  
for (j=0; j<M; j++)  
    p[j] = q[j] - r[j];  
printf("done\n");
```



AutoMutex: Critical Construct

```
sum = 0;
#pragma omp parallel private (lsum)
{
    lsum = 0;
    #pragma omp for
    for (i=0; i<N; i++) {
        lsum = lsum + A[i];
    }
    #pragma omp critical
    { sum += lsum; }
}
```

Threads wait their turn;
only one thread at a time
executes the critical section



Reduction Clause

Shared variable



```
sum = 0;  
#pragma omp parallel for reduction (+:sum)  
{  
    for (i=0; i<N; i++) {  
        sum = sum + A[i];  
    }  
}
```

OpenMP Schedule


- Can help OpenMP decide how to handle parallelism

`schedule(type [,chunk])`

- **Schedule Types**
 - **Static** – Iterations divided into size chunk, if specified, and statically assigned to threads
 - **Dynamic** – Iterations divided into size chunk, if specified, and dynamically scheduled among threads

Static Schedule


- Although the OpenMP standard does not specify how a loop should be partitioned
- Most compilers split the loop in N/p (N #iterations, p #threads) chunks by default.
- This is called a static schedule (with chunk size N/p)
 - *For example, suppose we have a loop with 1000 iterations and 4 omp threads. The loop is partitioned as follows:*



Static Schedule with chunk

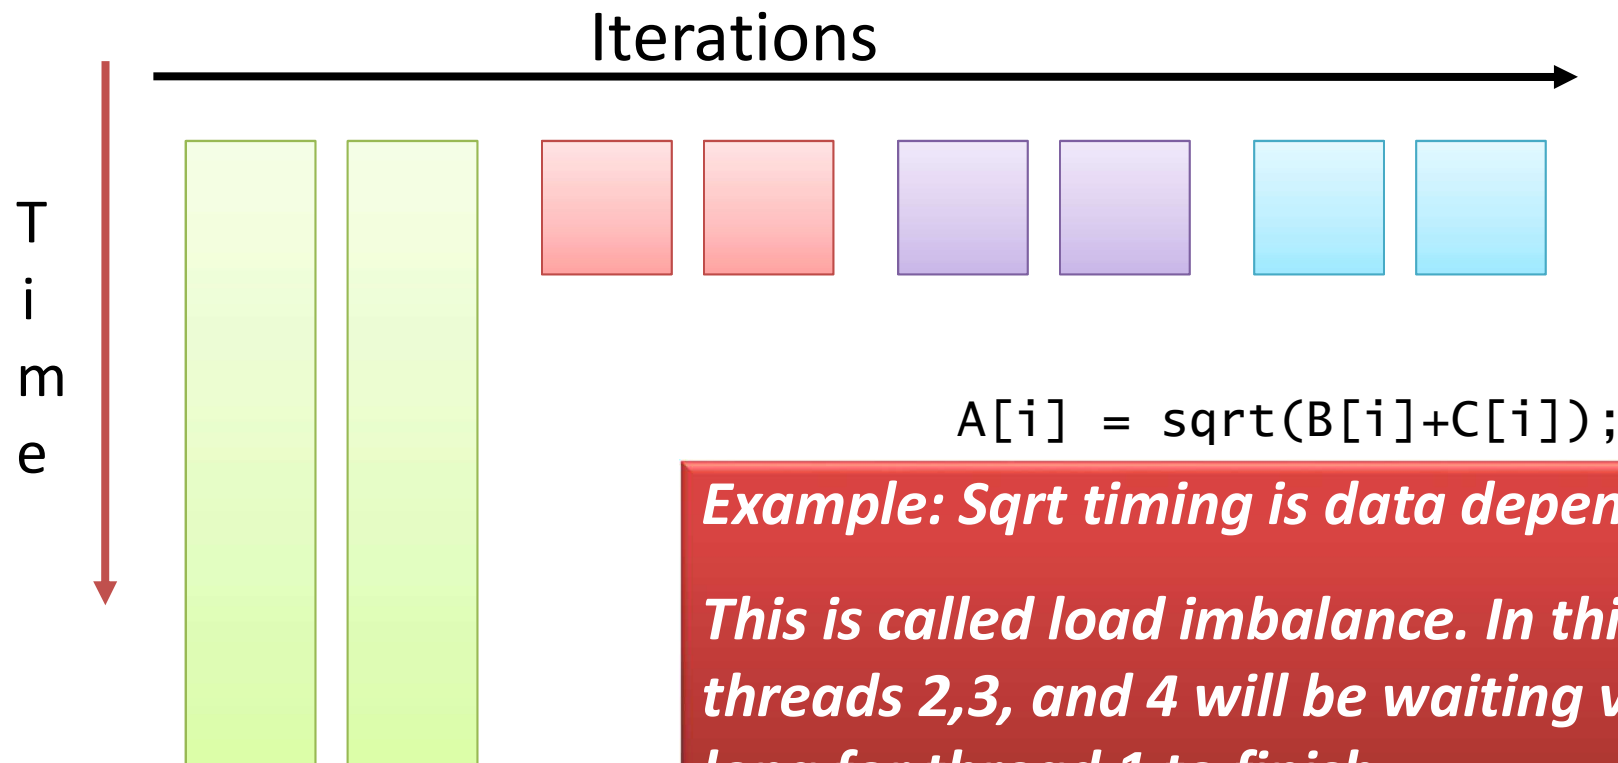
- A loop with 1000 iterations and 4 omp threads. Static Schedule with Chunk 10

```
#pragma omp parallel for schedule (static, 10)
{
  for (i=0; i<1000; i++)
    A[i] = B[i] + C[i];
}
```



Issues with Static schedule

- With static scheduling the number of iterations is evenly distributed among all openmp threads (i.e. Every thread will be assigned similar number of iterations).
- This is not always the best way to partition. Why is This?



Example: Sqrt timing is data dependent...

This is called load imbalance. In this case threads 2,3, and 4 will be waiting very long for thread 1 to finish

Dynamic Schedule

- With a dynamic schedule new chunks are assigned to threads when they come available.
- SCHEDULE(DYNAMIC,n)
 - Loop iterations are divided into pieces of size chunk. When a thread finishes one chunk, it is dynamically assigned another.

Dynamic Schedule

- SCHEDULE(GUIDED,n)
 - Similar to DYNAMIC but chunk size is relative to number of iterations left.
- Although Dynamic scheduling might be the preferred choice to prevent load imbalance
 - In some situations, there is a significant overhead involved compared to static scheduling.

More Examples on OpenMP

- <http://users.abo.fi/mats/PP2012/examples/OpenMP/>

Message Passing Interfaces

Outline

- Basic pf MPI
- MPI Constructs and Example
- Running programming in IITG HPC system
- Reference and Other Resources

How to compile and run on a Linux Machine

```
$sudo apt-get install openmpi*
```

```
$sudo apt-get install mpich*
```


```
$mpicc hello_mpi.c -o hello_mpi
```

```
$mpirun -np 4 ./hello_mpi
```

4 copies of hello_mpi process will run

The Message-Passing Programming Paradigm

- Message-Passing Programming Paradigm



The Message-Passing Programming Paradigm

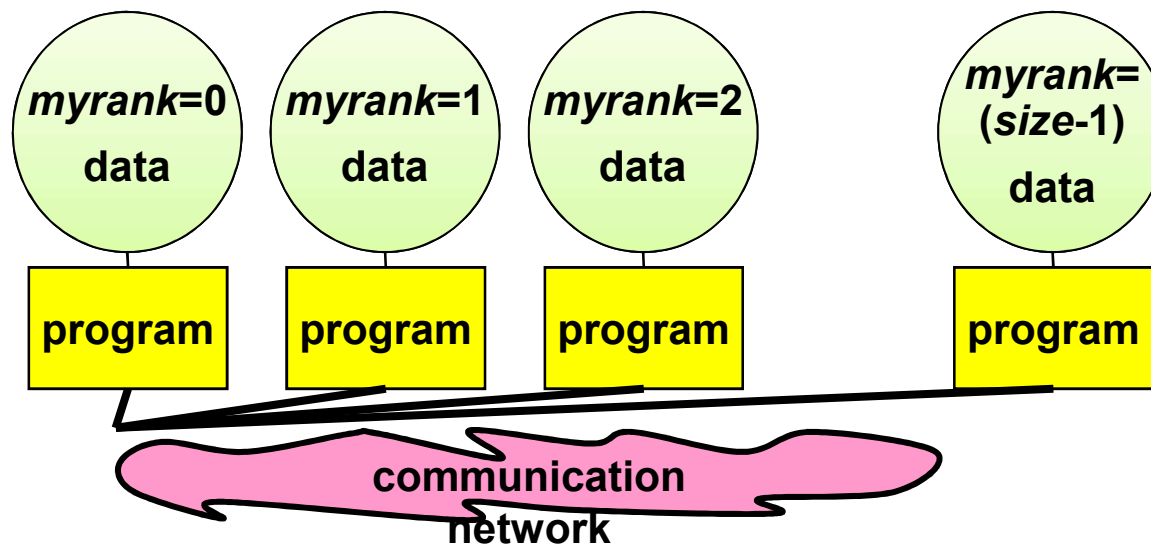
- Typically a single program operating on multiple datasets
- The variables of each sub-program have
 - The same name
 - But different locations (distributed memory) and different data!
 - i.e., all variables are local to a process
- Communicate via special send & receive routines (*message passing*)

Every process of MPI are different

- Hi : single person : you do
 - Touch you nose by left hand
 - Hi : Touch you head by right hand
- Hi: all persons of this hall do:
 - Touch your nose

Data and Work Distribution

- To communicate together mpi-processes need identifiers: **rank = identifying number**
- all distribution decisions are based on the *rank*
 - i.e., which process works on which data



What is SPMD


- **Single Program, Multiple Data**
- Same (sub-)program runs on each processor
- MPI allows also MPMD, i.e., **Multiple Program, ...**
 - but some vendors may be restricted to SPMD
 - MPMD can be emulated with SPMD

Emulation of MPMD

```
main(int argc, char **argv){  
    if (myrank < XX){  
        ocean( /* arguments */ );  
    }else{  
        weather( /* arguments */ );  
    }  
}
```

Message passing

- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
 - sending process
 - source location
 - source data type
 - source data size
 - receiving process i.e., the ranks
 - destination location
 - destination data type
 - destination buffer size



Access

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
 - phone line, mail box, fax machine, etc.
- MPI:
 - program must be linked with an MPI library
 - program must be started with the MPI startup tool

What is message passing?

- Data transfer
- Requires cooperation of sender and receiver
- Cooperation not always apparent in code

Blocking vs. Non-Blocking

- Blocking
 - The program will not continue
 - Until the communication is completed.
- Non-Blocking
 - The program will continue
 - Without waiting for the communication to be completed.

Features of MPI

- General
 - Communications combine context and group for message security.
 - Thread safety can't be assumed for MPI programs.

Features that are NOT part of MPI

- Process Management
- Remote memory transfer
- Threads
- Virtual shared memory

Why to use MPI?

- MPI provides a
 - Powerful, efficient, and portable way to express parallel programs.

Why to use MPI?

- MPI provides a
 - Powerful, efficient, and portable way to express parallel programs.
- MPI was explicitly designed
 - To enable libraries which eliminate the need for many users to learn **much inside of MPI.**

Why to use MPI?

- MPI provides a
 - Powerful, efficient, and portable way to express parallel programs.
- MPI was explicitly designed
 - To enable libraries which eliminate the need for many users to learn **much inside of MPI.**
- Portable !!!!!!!!!!!!!!!!!!!!!!!!!!!!!!!
- Good way to learn about subtle
 - Issues in parallel computing

How big is the MPI library?

- Huge (125 Functions).
- Basic (6 Functions).

How to install MPI in Linux Cluster

MPI Library

LAM : Local Area Multiprocessor

MPI CH: Argon National Laboratory

Skeleton MPI Program

```
#include <mpi.h>

void main( int argc, char **argv )
{
    MPI_Init( &argc, &argv );

    /* main part of the program */

    /*
    Use MPI function call depend on
    your data partitioning and the
    parallelization architecture
    */
    MPI_Finalize();
}
```

Initializing MPI

- The initialization routine MPI_INIT is the first MPI routine called.
- MPI_INIT is called once

```
int mpi_Init(  
    int *argc,  
    char **argv );
```


A minimal MPI program

```
#include "mpi.h"
#include <stdio.h>
int main(int argc, char *argv[])
{
    MPI_Init(&argc, &argv);
    printf("Hello, world!\n");
    MPI_Finalize();
    return 0;
}
```

How to compile and run on a Linux Machine

```
$mpicc hello_mpi.c -o hello_mpi
```

```
$mpirun -np 4 ./hello_mpi
```

**4 copies of hello_mpi process will
run**

A minimal MPI program cont.

- `#include <mpi.h>`
 - Provides basic MPI definitions and types.
- `MPI_Init` starts MPI
- `MPI_Finalize` exits MPI
- **Note that all non-MPI routines are local; thus “printf” run on each process**
- MPI functions return
 - Error codes or `MPI_SUCCESS`

Improved Hello.c

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char *argv[])
{
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    printf("I am %d of %d\n", rank, size);
    MPI_Finalize();
    return 0;
}
```

MPI Basic Communication

Constructs

MPI Concepts

- The default communicator is the
`MPI_COMM_WORLD`
- A process is identified
 - By its rank in the group associated with a communicator.

Finding Out About the Environment

- Two important questions that arise early in a parallel program are:
 - How many processes are participating in this computation?
 - Which one am I?

Finding Out About the Environment

- MPI provides functions to answer these questions:
 - **MPI_Comm_size** reports the number of processes.
 - **MPI_Comm_rank** reports the *rank*, a number between 0 and size-1, identifying the calling process

MPI: Data Types

- The data message which is sent or received is described by a triple
 - Address, count, data type
- The following data types are supported
 - Predefined data types
 - Arrays, sub blocks of a matrix and user defined

MPI blocking send

```
MPI_Send(void *start, int count,  
         MPI_DATATYPE datatype,  
         int Dest, int Tag,  
         MPI_COMM comm );
```

MPI blocking send

```
MPI_Send(void *start, int count,  
         MPI_DATATYPE datatype,  
         int Dest, int Tag,  
         MPI_COMM comm );
```

- The message buffer is described by (start, count, datatype).
- Dest is the rank of the target process in the defined communicator
- Tag is the message identification number.

MPI blocking receive

```
MPI_Recv(void *start, int count,  
         MPI_DATATYPE datatype,  
         int Source, int Tag,  
         MPI_COMM comm, MPI_COMM *S) ;
```

MPI blocking receive

```
MPI_Recv(void *start, int count,  
         MPI_DATATYPE datatype,  
         int Source, int Tag,  
         MPI_COMM comm, MPI_COMM *S) ;
```

- **Source** is the rank of the sender in the communicator.
- The receiver can specify
 - A wildcard value for source (MPI_ANY_SOURCE)
 - A wildcard value for tag (MPI_ANY_TAG),
 - Indicating that any source and/or tag are acceptable


MPI blocking receive

```
MPI_Recv(void *start, int count,  
         MPI_DATATYPE datatype,  
         int Source, int Tag,  
         MPI_COMM comm, MPI_COMM *S) ;
```

- **Status** is used for extra information
 - About the received message if a wildcard receive mode is used
- If the count of the message received is \leq described by the MPI receive command
 - Message is successfully received
 - Else it is considered as a buffer overflow error.

Sources of Deadlocks

- Send a large message from proc 0 to proc 1
 - If there is insufficient storage at the destination, the send must wait for the user to provide the memory space (through a receive)
- What happens with



- This is called “unsafe” because it depends on the availability of system buffers

MPI Collective Communication


Constructs

Collective Communication

- **Optimised Communication routines involving a group of processes**
- Collective action over a communicator
 - i.e. all processes must call the collective routine
- Synchronization may or may not occur
- All collective operations are blocking
 - and No tags
- Receive buffers must have exactly the same size as send buffers.

Broadcast

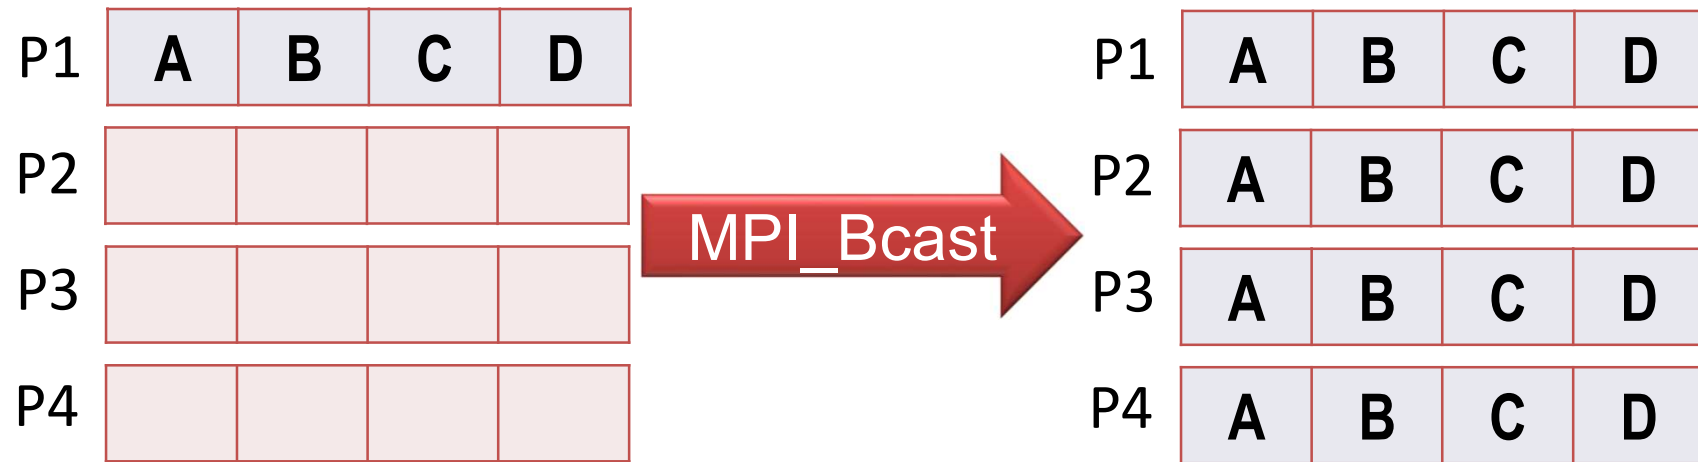
```
int MPI_Bcast(void *buf, int count,  
             MPI_Datatype datatype,  
             int root, MPI_Comm comm);
```



e.g., root=1

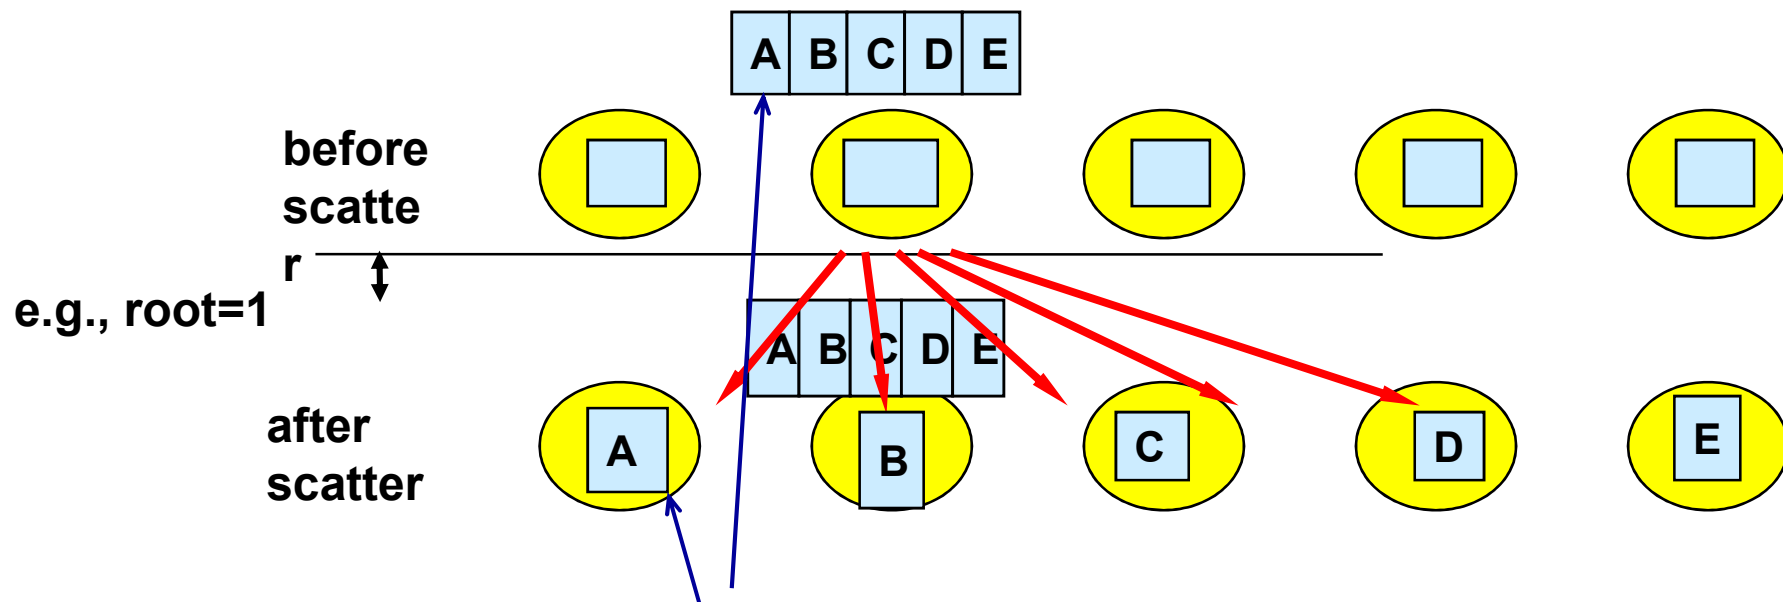
- rank of the sending process (i.e., root process)
- must be given identically by all processes

Broadcast



Scatter

```
int MPI_Scatter(void *sendbuf, int  
    sendcnt, MPI_Datatype sendtype,  
    void *recvbuf, int recvcount,  
    MPI_Datatype recvtype,  
    int root, MPI_Comm comm);
```



Scatter


P1	A	B	C	D
P2				
P3				
P4				

MPI_Scatter


P1	A	B	C	D
P2	B			
P3	C			
P4	D			

Gather


```
int MPI_Gather(void *sendbuf,  
              int sendcnt, MPI_Datatype  
              sendtype, void *recvbuf,  
              int recvcnt, MPI_Datatype  
              recvtype, int root, MPI_Comm comm)
```



Gather



Gather to All



Global Reduction Operations

- Perform a global reduce operation across all members of a group.


$$d_0 \circ d_1 \circ d_2 \circ d_3 \circ \dots \circ d_{s-2} \circ d_{s-1}$$

- d_i = data in process rank i
 - single variable, or vector
- \circ = associative operation
- Example:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation

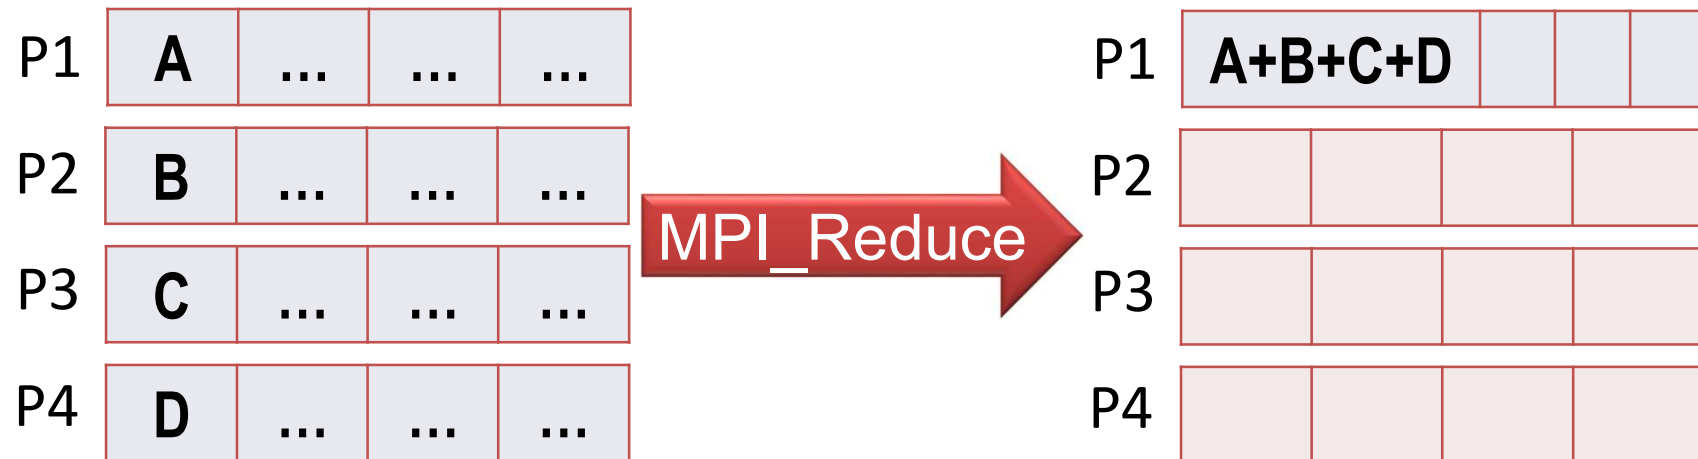
MPI Reduce

before MPI_REDUCE

- inbuf and result




MPI_Reduce



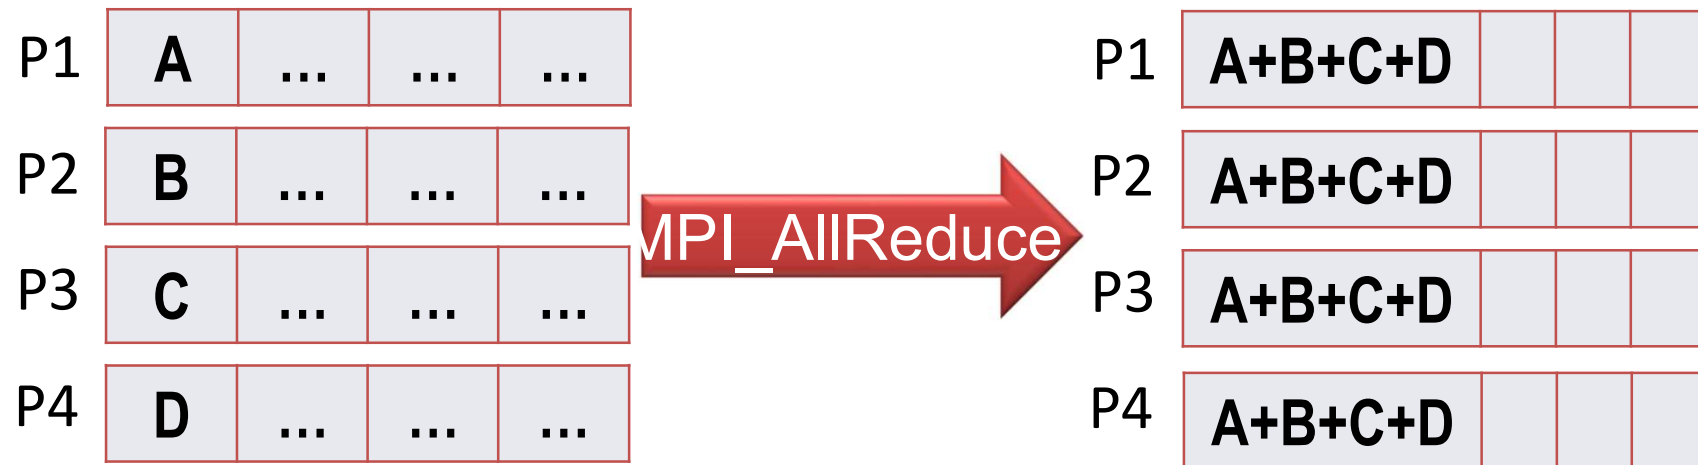
MPI_AllReduce

before MPI_ALLREDUCE

- inbuf and result



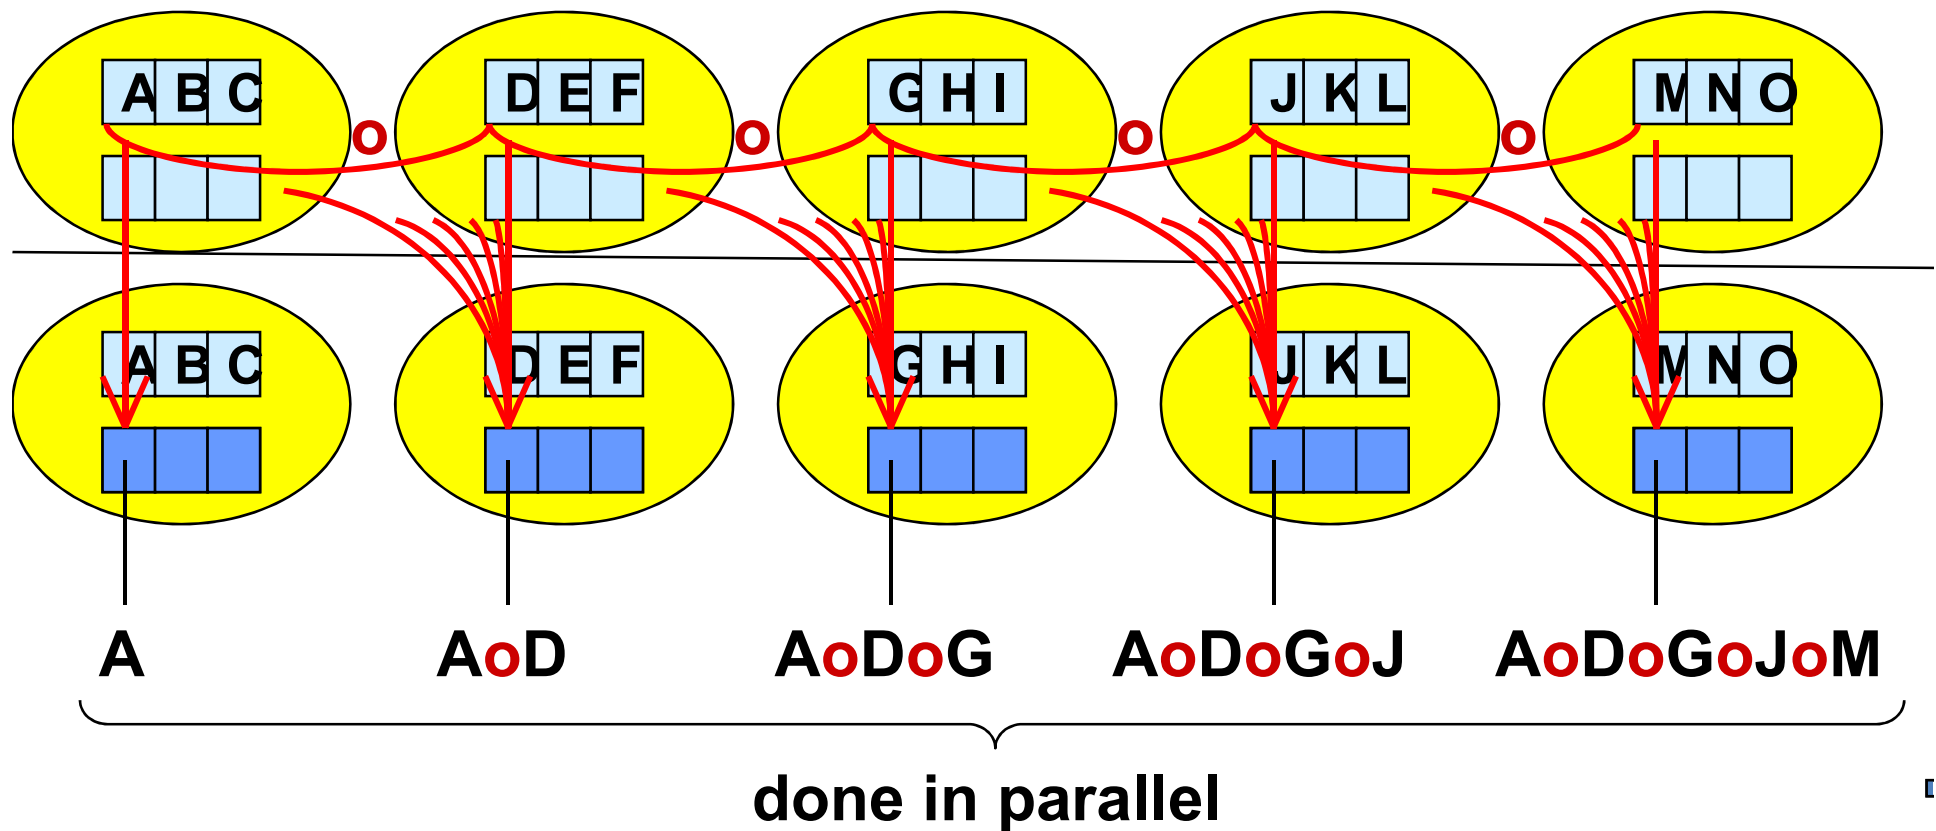
MPI_AllReduce



MPI_Scan

before MPI_SCAN

- inbuf and result



MPI Examples


Example: Sum of N data

- Master Process
 - Data to be read by process 0 or MASTER
 - Divide the data in to N/M chunk size ($N \% M == 0$)
 - **SEND** respective chunk of data to other process
 - **Do local sum on each process (in master also)**
 - **RECV** sum of other process and calculate final sum
- Other Process
 - **RECV** data from Mater
 - **Do local sum on each process**
 - **SEND** local sum to MASTER

See the Code

Example: Matrix MUL

- $c = a \times b$: $a[NRA][NCA]$, $b[NCA][NCB]$, $c[NRA][NCB]$
- Work get divided: Based on Rows



Example: Matrix MUL

- $c = a \times b$: $a[NRA][NCA]$, $b[NCA][NCB]$, $c[NRA][NCB]$
- One Master Processor
- Many Workers, Assume $NRA \% NumWorker == 0$
 - Master divide the work between worker
 - Send respective rows of A and whole B to workers
 - RECV array C from all worker
- Every Worker
 - get some Row of A, Whole of B
 - calculate part of C
 - Send calculated C to Master

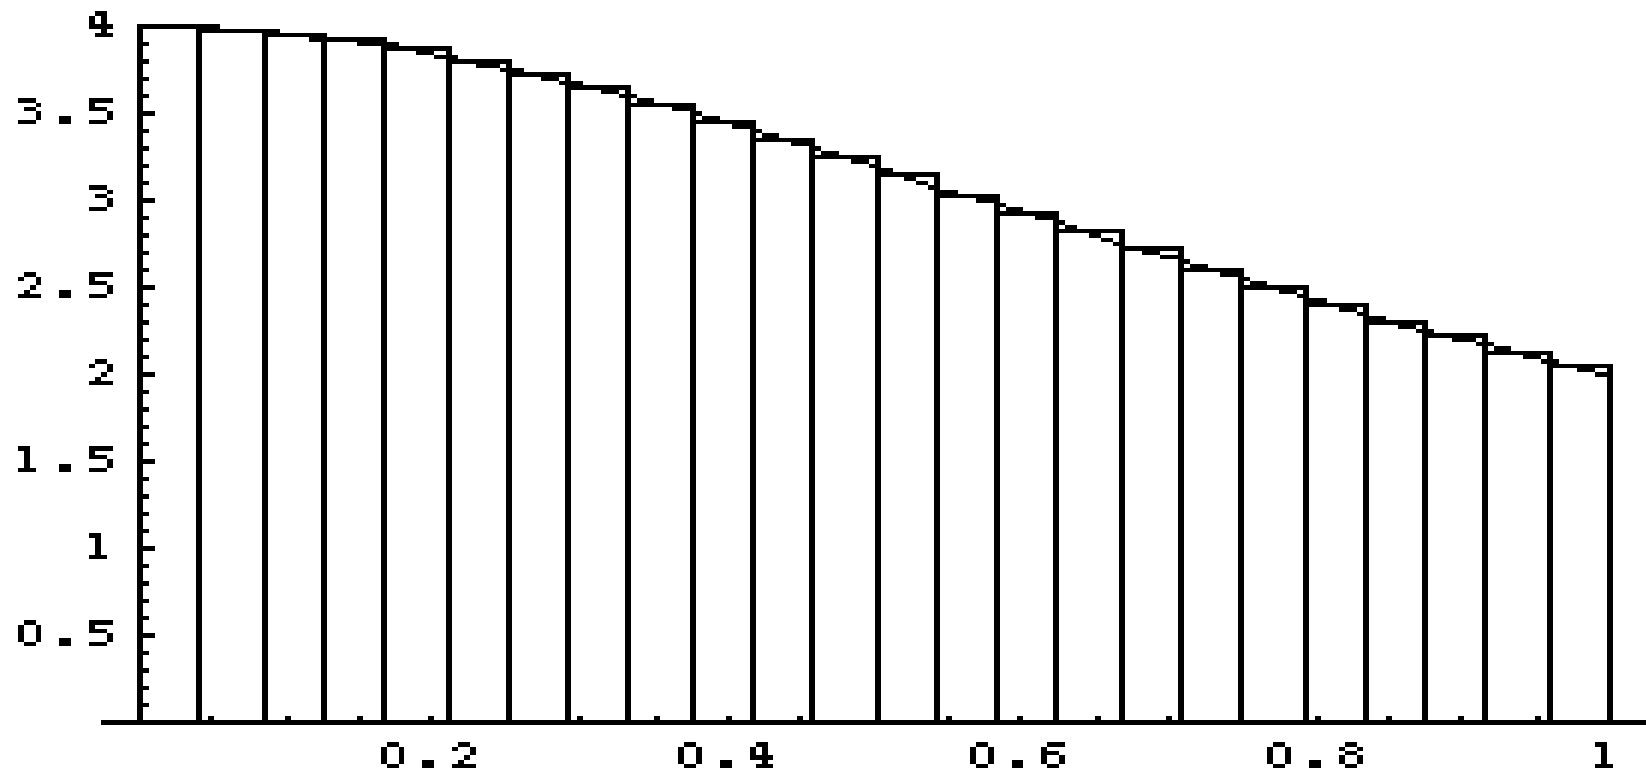
See the Code

Example: Compute PI

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$

Example: Compute PI

$$\pi = \int_0^1 \frac{4}{1+x^2} dx$$



How to write Program?


- Divide the range in to N interval/piece
 - Piece of size $h = \text{Range}/N$;
- Calculate area under each piece
 - Calculate the function value at piece X and multiply with piece size
 - $h * F(X)$
- Sum all the piece
 - $\sum_{i=1}^n h * F(X_i)$ with $X_i = R_{\min} + i * h$

How to write Program?

```
printf("Enter Num intervals: ");
scanf("%d", &n);
h = 1.0 / (double)n;
sum = 0.0;
for (i=1; i<n; i++) {
    x = h*(i-0.5); Fx=4.0/(1.0+ x*x);
    sum = sum + Fx;
}
pi = h*sum;
printf("pi is approx %.16f", pi);
```



How to write Parallel Program?

- Divide the range in to N interval/piece
 - Piece of size $h = \text{Range}/N$;
 - **Suppose $N = 1000$, NumProcessor = 4**
- In Parallel: Calculate area under each piece



How to write Parallel Program?

- Divide the range in to N interval/piece
 - Piece of size $h = \text{Range}/N$;
 - **Suppose $N = 1000$, NumProcessor = 4**
- In Parallel: Calculate area under each piece



- $(hF_{x1} + hF_{x5} + \dots + hF_{x997}) + (hF_{x2} + hF_{x6} + \dots + hF_{x998}) +$
 $(hF_{x3} + hF_{x7} + \dots + hF_{x999}) + (hF_{x4} + hF_{x8} + \dots + hF_{x996})$

Example: Compute PI

```
#include <mpi.h>
#include <math.h>
int main(int argc, char *argv[]) {
    int n, myid, Nproc, i;
    double lsum, pi, h, sum, x, a;
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &Nproc);
    MPI_Comm_rank(MPI_COMM_WORLD, &myid);
    if (myid == 0) {
        printf("Enter Num intervals: \n");
        scanf("%d", &n);
    }
    MPI_Bcast(&n, 1, MPI_INT, 0,
              MPI_COMM_WORLD);
```

Example: Compute PI

```
h = 1.0 / (double)n; sum = 0.0;
for (i=myid+1; i<=n; i+= Nproc) {
    x = h*((double)i - 0.5);
    sum += 4.0 / (1.0 + x * x);
}
lsum = h*sum;
MPI_Reduce(&lsum, &pi, 1, MPI_DOUBLE,
           MPI_SUM, 0, MPI_COMM_WORLD);
if (myid == 0)
    printf("pi is approx %.16f\n", pi);
MPI_Finalize();
return 0;
}
```

IITG HPC clusters: Spec

- 4 login nodes
- 126 compute node
- 16 GPU compute nodes
- 16 Phi compute nodes
- Total $126+16+16= 158$ nodes
 - Each node 12 cores * 2 threaded
 - Effective $24*158 = 3792$ cores

Running MPI program on IITG HPC clusters

- Logic to one login nodes : non GPU/PHI
 - param.-ishan.iitg.ernet.in (172.17.0.7)
- Compile MPI-code

Running MPI program on IITG HPC clusters

- Logic to one login nodes : non GPU/PHI
 - param.-ishan.iitg.ernet.in (172.17.0.7)
- Compile MPI-code
- Run using srun or sbatch
 - In s batch specify number of node, task per node
 - Total process
- SLURM : Simple Linux Util for Resce Mngt
 - Scheduler the JOB efficiently, user need not to worry where it is scheduling

Resources

- <https://computing.llnl.gov/tutorials/mpi/>
- V. Kumar, A. Grama, A. Gupta, and G. Karypis. ***Introduction to Parallel Computing: Design and Analysis of Algorithms***. Benjamin-Cummings Publ. Co, 1994 [metis software]
- Michael J. Quinn. ***Parallel Programming in C with MPI and OpenMP***. McGraw-Hill Education Group. 2003.
- Joseph JáJá. ***An Introduction to Parallel Algorithms***. Addison Wesley Longman Publishing Co., Inc.,, USA. 1992